

IMPLEMENTATION OF DATABASE SYSTEMS

Fall 2021

Instructor:	Amir Hossein Rouhani Seraji	Code:	13755
Email:	ah.rouhaniseraji@gmail.com	Unit:	3

Course Pages: Please check this website two or three times a week

- <https://rouhani-class.github.io/idbms001>

Discussion Group: Every student has to be a member of our Telegram Channel and Telegram Group for our discussion and daily announcements.

- [Telegram Channel](#).
- [Telegram Group](#).

Teaching Assistant:

- [Mehran Moeini Jam](#).
- [Mohammad Hossein Khoshechin](#).

Prerequisites: An undergraduate-level understanding of Relational Databases Concepts, Data Structures, and Algorithms Design & Analysis.

Objectives: This is an upper-level course on the internals of database management systems. This course is on the design and implementation of database management systems. Topics include storage models (n-ary, decomposition), storage architectures (heaps, log-structured), indexing (order preserving trees, hash tables), transaction processing (ACID, concurrency control), recovery (logging, checkpoints), query processing (joins, sorting, aggregation, optimization), and parallel architectures (multi-core, distributed). Upon successful completion of this course, the student should be able to: Implement a disk-oriented database storage manager with table heaps and indexes. Understand, compare, and implement the fundamental concurrency control algorithms. Implement database recovery algorithms and verify their correctness. Identify trade-offs among database systems techniques and contrast distributed/parallel alternatives for both on-line transaction processing and on-line analytical workloads. Interpret and comparatively criticize database system architectures.

- Introduction
 - Review Relational Theory Concepts
 - Review Relational Algebra
 - Review SQL Languages
 - History of Data Models and Database Systems
- Physical Storage Systems
 - Overview of Physical Storage Media
 - Storage Interfaces
 - Magnetic Disks

- Flash Memory
- Data Storage Structures
 - File Organization
 - Organization of Records in Files
 - Data-Dictionary Storage
 - Database Buffer
- Indexing
 - Ordered Indices
 - B⁺-Tree Index Files
 - B⁺-Tree Extensions
 - Hash Indices
- Query Processing
 - Measures of Query Cost
 - Selection Operation
 - Sorting
 - Join Operation
- Query Optimization
 - Transformation of RelationalExpressions
 - Estimating Statistics of ExpressionResults
 - Choice of Evaluation Plans
 - Materialized Views
 - Advanced Topics in QueryOptimization
- Transactions
 - Transaction Concept
 - Transaction Atomicity and Durability
 - Transaction Isolation
 - Serializability
 - Implementation of Isolation Levels
- Concurrency Control
 - Lock-Based Protocols
 - Deadlock Handling
 - Multiple Granularity
 - Timestamp-Based Protocols
 - Validation-Based Protocols
 - Multiversion Schemes
 - Snapshot Isolation
 - Weak Levels of Consistency inPractice

- Recovery from deadlock
- Recovery System
 - Failure Classification
 - Recovery and Atomicity
 - Recovery Algorithms
 - Buffer Management
 - Failure with Loss of Non-Volatile Storage
 - High Availability Using Remote Backup Systems
 - Early Lock Release and Logical Undo Operations
 - ARIES
- Database-System Architectures
 - Parallel Systems
 - Distributed Systems
 - Cloud-Based Services
- Parallel and Distributed Storage
 - Data Partitioning
 - Dealing with Skew in Partitioning
 - Replication
 - Parallel Indexing
 - Distributed File Systems
 - Parallel Key-Value Stores
- Parallel and Distributed Query Processing
 - Parallel Sort
 - Parallel Join
 - Parallel Evaluation of Query Plans
 - Query Processing on Shared-Memory Architectures
 - Query Optimization for Parallel Execution
 - Parallel Processing of Streaming Data
 - Distributed Query Processing
- Parallel and Distributed Transaction Processing
 - Distributed Transactions
 - Commit Protocols
 - Concurrency Control in Distributed Databases

Main References: Our Lectuers are mainly based-on this book.

- Abraham Silberschatz, Henry F. Korth, S. Sudarshan, *Database System Concepts*, McGraw-Hill, 7th ed, 2020.

Supplementary References: We also use these references for complementary discussion.

- Goetz Graefe, *On Transactional Concurrency Control*, Morgan & Claypool, 2019.
- M. Tamer Ozsu, Patrick Valduriez, *Principles of Distributed Database Systems*, Springer, 4th ed, 2020.
- Gerhard Weikum, Gottfried Vossen, *Transactional Information Systems Theory, Algorithms, and the Practice of Concurrency Control and Recovery*, Morgan Kaufmann, 2001.
- Jim Gray, Andreas Reuter, *Transaction Processing: Concepts and Techniques*, Morgan Kaufmann, 1993.
- Philip A. Bernstein, Eric Newcomer, *Principles of Transaction Processing*, Morgan Kaufmann, 2nd ed, 2009.
- Sam S. Lightstone, Toby J. Teorey, Tom Nadeau, *Physical Database Design: The Database Professional's Guide to Exploiting Indexes, Views, Storage, and More*, Morgan Kaufmann, 4th ed, 2007.
- Nicolas Bruno, *Automated Physical Database Design and Tuning*, CRC Press, 2011.

Useful Links: Courses with Video Lectures.

- [CMU 15-445/645 Database Systems](#), Carnegie Mellon University, Fall 2021.
- [CMU 15-445/645 Database Systems](#), Carnegie Mellon University, Fall 2019.
- [CMU 15-445/645 Database Systems](#), Carnegie Mellon University, Fall 2018.
- [CMU 15-445/645 Database Systems](#), Carnegie Mellon University, Fall 2017.
- [CMU 15-721 Advanced Database Systems](#), Carnegie Mellon University, Spring 2020.
- [CMU 15-721 Advanced Database Systems](#), Carnegie Mellon University, Spring 2019.
- [CMU 15-721 Advanced Database Systems](#), Carnegie Mellon University, Spring 2018.
- [CMU 15-721 Advanced Database Systems](#), Carnegie Mellon University, Spring 2017.
- [CMU 15-721 Advanced Database Systems](#), Carnegie Mellon University, Spring 2016.

Seminars.

- [Vaccination Database Tech Talks Second Dose](#) - 2021.
- [Vaccination Database Tech Talks First Dose](#) - 2021.
- [Quarantine Database Tech Talks](#) - 2020.
- [Hardware Accelerated Database Lectures](#) - Fall 2018.
- [Time Series Database Lectures](#) - Fall 2017.
- [The Databaseology Lectures](#) - Fall 2015.
- [Seven Databases in Seven Weeks](#) - Fall 2014.
- [A Celebration of Mike Stonebraker](#) - 2014.

Research Groups in Database Systems.

- [Carnegie Mellon Database Group](#).
- [MIT Database Group](#).
- [MIT Data Systems Group](#).
- [UC Berkeley Data Systems and Foundations](#).
- [Stanford Infolab Group](#).
- [University of Washington's database group](#).
- [UW-Madison Database Group](#).
- [Database of Databases](#)

Books website.

- [Database System Concepts](#).

Our passed Semesters Audio Lectures.

- [7232 Principles of Database Design, IAU South Tehran Branch, Fall 2018](#).
- [7232 Principles of Database Design, IAU South Tehran Branch, Spring 2017](#).

Grading Policy:

- Assignments (40%)
- Quiz (30%)
- Final Exam (30%)
- (Optional) Presentation (10%)